

The Accuracy Illusion: How Overlapping Samples Manufacture 71 % Hit-Rates in Intraday Index Forecasting

Rachit Ranka

Department of Computer Science Engineering
Chandigarh University
Mohali, India
rankarachit5@gmail.com

Aaditya Saini

Department of Computer Science Engineering
Chandigarh University
Mohali, India
aadityasaini899@gmail.com

Abstract—Intraday forecasting systems routinely report directional hit-rates between 60 % and 85 %, and those figures move real capital. This paper shows that they are not valid estimates of predictive skill, and measures precisely how they arise. When a k -bar-ahead forecast is issued every bar, consecutive forecasts share $k-1$ of their k outcome bars. If the model’s own calls also persist, and on a live production system we measure a mean run of 15.2 minutes, the reported hit-rate acquires a variance far larger than its nominal sample size implies. We simulate a forecaster with *exactly zero skill* against real NSE index bars. It posts a ≥ 60 % session once in six and a ≥ 71 % session once in 51, while reporting a confidence interval $1.66\times$ too narrow. Across a panel of two indices and four horizons, a ≥ 71 % reading surfaces every seventh session. Under correct per-event scoring, nine model families spanning logistic regression to a Temporal Fusion Transformer, all trained on identical purged walk-forward splits, reach at most 51.3 % out-of-sample. A sweep of 86 rule-based strategies yields a best Sharpe *below* what a circular block-bootstrap null produces from selection alone ($P = 0.658$). We give the closed form behind the distortion, showing the variance inflation is governed by the second moment of the run-length distribution, and an estimator built on it lifts the coverage of a nominal 95 % interval from 40 % to 94 %; the corrected effective sample size of a 69-forecast session is six. We then quantify what would actually be required. Break-even accuracy derived from measured move distributions is 71.5 % at the 30-minute horizon, and it exceeds 100 % at shorter horizons and at option-level costs. The one directional effect that survives every control (+1.6 points over a shuffled-label baseline, $p = 0.001$, Bonferroni-robust) is therefore genuine and an order of magnitude too small to trade. All code, raw results and the free-data corpus are public.

Index Terms—overlapping observations, backtest overfitting, directional forecasting, market microstructure, multiple hypothesis testing, negative results, reproducibility, index derivatives

I. INTRODUCTION

A forecasting system that reports 71 % directional accuracy is, on its face, extraordinary. The literature places realistic out-of-sample directional accuracy for liquid index

intraday prediction near the coinflip, and treats figures above 60 % with suspicion [16], [27]. Yet such numbers appear regularly on trading dashboards, in vendor marketing and in applied publications, and they are frequently computed in good faith from data the system itself generated.

This paper began as an attempt to build such a system. Over fourteen months we constructed a live multi-horizon forecasting platform for Indian index derivatives, and it reported a headline accuracy of 71 %. That number was wrong. Tracing it to source revealed no bug in the model. The defect lay in the *denominator*, in how a per-minute forecast stream is scored. Four further positive results from the same project were later retracted for related reasons.

The econometrics of overlapping observations is not new [1], [2]. Hansen and Hodrick [1] established four decades ago that sampling data more finely than the forecast interval induces serial correlation in forecast errors and invalidates naive standard errors. What that literature does not document, and what we supply here, is the *magnitude and shape* of the distortion in a now-ubiquitous practice: the every-minute multi-horizon directional dashboard. There the overlap meets a second ingredient, prediction persistence, and the product is then amplified by the multiplicity of an index $\times$ horizon panel.

A. Contributions

- 1) **A measurement of the artefact as it occurs in practice.** Overlap on its own changes nothing, since a forecaster that re-draws every bar yields an i.i.d. hit indicator. Inflation requires overlap *and* persistent predictions together. We calibrate persistence from a production system’s own forecast log, then quantify the resulting distribution, its calibration error and its panel multiplicity (Section V).
- 2) **A complete control set for a negative result.** Shuffled-label controls on every model, injected synthetic signals of known strength as a positive control, nine architectures under identical purged splits, and a block-bootstrap null for the strategy sweep (Section VII). The positive control is essential. Without

it, “we found nothing” cannot be distinguished from “our pipeline is broken.”

- 3) **A closed form and a corrected estimator.** We derive the effective sample size of an overlapping, persistent forecast stream and show it is governed by $\mathbb{E}[L^2]/\mathbb{E}[L]$ over the run lengths. An interval built on it covers at 94% where the naive interval covers at 40% and a standard HAC correction at 67% (Section VI).
- 4) **The required-versus-achieved framing.** We derive break-even accuracy from measured move distributions and show that below a horizon threshold *no* accuracy is sufficient, which reframes high-accuracy claims as either artefacts or as far less valuable than they appear (Section VIII).
- 5) **A reusable free-data corpus** for NSE index derivatives, including a decade of end-of-day option chains and two years of participant-wise open interest, with the observation that the exchange’s static archive is not geo-restricted while its live API is (Section III).

All experiments are reproducible from public code and committed raw results.¹

II. RELATED WORK

A. Overlapping observations

Hansen and Hodrick [1] and Richardson and Smith [2] characterise inference under overlapping returns; Newey and West [3] give the standard heteroskedasticity- and autocorrelation-consistent correction. That literature concerns regression coefficients and their standard errors. We apply the same logic to a quantity practitioners actually report, the session directional hit-rate, and measure the consequence empirically rather than asymptotically.

B. Backtest overfitting and multiple testing

Harvey, Liu and Zhu [10] argue that the conventional $t > 2$ threshold is indefensible after decades of factor mining and propose $t > 3$. Bailey and López de Prado [11] introduce the Deflated Sharpe Ratio, and Bailey *et al.* [12] the Probability of Backtest Overfitting via combinatorially symmetric cross-validation. White’s Reality Check [9] and the Benjamini–Hochberg procedure [8] provide complementary family-wise and false-discovery-rate controls. We apply all of these, and report openly where each one fails on our data and why.

C. Machine learning for financial prediction

Fischer and Krauss [25] apply LSTM networks to S&P 500 constituents, reporting daily returns of 0.46% and a Sharpe ratio of 5.8 *before transaction costs*; Krauss *et al.* [26] compare deep networks, gradient-boosted trees and random forests in the same setting. Gu, Kelly and Xiu [24] conduct a comprehensive machine-learning comparison for

asset pricing, reporting modest out-of-sample R^2 . Zhang *et al.* [23] achieve strong accuracy on limit-order-book data, and Kolm *et al.* [22] demonstrate genuine multi-horizon alpha from order-flow imbalance. Both results rest on message-level exchange data rather than on bars, which is the distinction that matters for our purposes. Sezer *et al.* [27] survey the field.

Table I compares this work with representative prior studies along the methodological axes that determine whether a reported accuracy is interpretable. We deliberately compare on *procedure* rather than on headline numbers: our claim is not that prior results are wrong, but that several widely used reporting practices admit the artefact documented here, and that a reported accuracy cannot be assessed without knowing the scoring unit and the cost model.

D. What is predictable

Volatility clustering [17], [18] and the intraday volatility U-shape [19] are among the most robust regularities in financial time series, and they concern magnitude rather than direction. The variance risk premium, meaning the persistent gap between implied and subsequently realised volatility, is documented by Carr and Wu [20] and by Bakshi and Kapadia [21]. Our positive results fall entirely within these two categories, and we regard that as the paper’s substantive finding rather than as a consolation.

III. DATA

All data is obtainable free of charge from a retail brokerage account and public exchange archives. Table II summarises the corpus.

Three observations about this corpus are of independent practical value.

Index spot carries no volume. The cash index prints zero volume, so any feature computed as a “VWAP” on index spot is in fact an unweighted expanding mean of typical price. This was live in our own feature set for weeks before it was identified. Real volume exists only on the futures contract.

The static exchange archive is not geo-restricted. The exchange’s *live* option-chain API returns HTTP 403/404 from outside India, but the static end-of-day derivatives archive downloads normally. This distinction converted an assumed hard constraint into a two-year, and subsequently ten-year, dataset.

Time-based one-time-password authentication is location-independent. Brokers authenticating by locally computed TOTP permit data collection from any jurisdiction, whereas SMS-based flows do not. This is a deciding practical constraint on reproducibility, and it is not documented prominently anywhere.

IV. METHODOLOGY

Every accuracy reported in this paper passes through one shared harness, so that differences between models are differences between models.

¹<https://github.com/rachitranka25/accuracy-illusion>

TABLE I

METHODOLOGICAL COMPARISON WITH REPRESENTATIVE PRIOR WORK. “SCORING UNIT” IS THE UNIT IN WHICH HIT-RATE OR RETURN IS AGGREGATED; PER-BAR SCORING OF A MULTI-BAR HORIZON IS THE CONDITION UNDER WHICH THE ARTEFACT OF SECTION V APPEARS. MT: MULTIPLE-TESTING CORRECTION. NC: NULL CONTROL (SHUFFLED LABELS OR EQUIVALENT). C: TRANSACTION COSTS MODELLED. R: CODE AND RAW RESULTS PUBLIC.

Study	Market	Data granularity	Scoring unit	MT	NC	C	R
Fischer & Krauss [25]	S&P 500	Daily bars	Per day, non-overlapping	—	—	pre-cost	—
Krauss <i>et al.</i> [26]	S&P 500	Daily bars	Per day, non-overlapping	—	—	pre-cost	—
Gu <i>et al.</i> [24]	US equities	Monthly	Per month	partial	—	partial	partial
Zhang <i>et al.</i> [23]	LSE	Message-level LOB	Per event	—	—	—	partial
Kolm <i>et al.</i> [22]	Nasdaq	Message-level LOB	Per horizon	—	—	discussed	—
Applied index studies (typical)	NSE, global	Daily / intraday bars	Often per bar	—	—	—	rare
This work	NSE index	5-min bars, EOD chains	Per event (k-bar)	✓	✓	✓	✓

This work applies Bonferroni, Benjamini–Hochberg, the Deflated Sharpe Ratio and PBO/CSCV (MT); shuffled-label controls on every model plus injected positive controls (NC); explicit round-trip costs, and inverts them into a required-accuracy threshold (C).

TABLE II
FREE-DATA CORPUS. ALL COLLECTORS ARE PUBLIC.

Asset	Period	Size
BANKNIFTY 5-min bars	2025-02 – 2026-07	26 265
NIFTY 50 5-min bars	2025-02 – 2026-07	26 267
3 further indices, 5-min	2025 – 2026	9.6k–14k
Option chain EOD + OI	2024-07 – 2026-07	1 382 738
Long-history option chain	2016-05 – 2026-05	2 181 d
Participant-wise OI	2024-07 – 2026-07	515 d
Index futures (real volume)	2024-01 – 2026-07	639 d

Purged walk-forward. Expanding-window splits, with the k bars whose labels overlap the test window removed from the end of each training set. Label overlap is the most common leakage channel in this setting [13].

Exact intervals. Accuracies are reported with Clopper–Pearson 95% confidence intervals and the exact binomial p -value against 0.5. We additionally report the Pesaran–Timmermann statistic [6], which is the appropriate null for directional accuracy: it compares the realised hit-rate against what independence between forecast and outcome would produce given their own marginal rates, and therefore catches a model that merely tracks drift.

Shuffled-label controls. Every model is additionally trained and evaluated with randomly permuted labels. If real and shuffled scores coincide, the pipeline is reading noise. That is a stronger statement than any train/test gap, since it does not require the model to be well specified.

Positive controls. A synthetic feature agreeing with the label on a known fraction s of rows is injected, making the recoverable accuracy approximately $0.5 + s/2$. A pipeline that fails to recover it cannot be trusted to report a null.

Per-event scoring and costs. Strategy returns are evaluated on non-overlapping k -bar holds with explicit round-trip costs, and costs are treated as part of the hypothesis rather than as a subsequent adjustment.

Features are the fifteen columns used by the production system: one-bar return, RSI, EMA distance, MACD

histogram, Bollinger width and %b, ATR%, Parkinson volatility, VWAP distance, time-of-day, day-of-week, regime volatility, Hurst exponent, and distance from the previous session’s high and low. All are transforms of a single price series, a point we return to in Section IX.

V. THE MECHANISM

A. Overlap is arithmetic, not information

Fig. 1 shows the autocorrelation of the k -bar forward return measured two ways. Computed on overlapping windows it appears strongly autocorrelated, at 0.841 for lag 1, and decays to approximately zero at *exactly* lag k . Sampled so that no two observations share a bar, the autocorrelation is 0.067 at lag 1 and indistinguishable from zero thereafter. The pattern reproduces independently at $k = 6$ (30 minutes) and $k = 12$ (60 minutes), with the decay landing on lag 6 and lag 12 respectively. The correlation is a property of the sampling scheme, not of the market.

B. Persistence is the second, necessary ingredient

Overlap alone is insufficient. If predictions are drawn independently each bar, the hit indicator $\mathbf{1}\{\hat{y}_t = y_t\}$ is i.i.d. Bernoulli regardless of how autocorrelated y_t is, and no inflation occurs. Our first simulation made precisely this error and correctly found nothing.

Real models do not re-decide every bar: their features move slowly, so they repeat a call for long stretches. Measured on the production system’s own forecast log at the 30-minute horizon, BANKNIFTY calls persist for a mean of 15.2 minutes (99 runs across 1 500 forecasts, longest 135 minutes) and NIFTY 50 for 12.2 minutes (123 runs, longest 207 minutes). A model re-drawing each bar would show a mean run of one.

When both the call and the outcome persist, a single trending stretch is recorded as dozens of hits.

C. What a zero-skill forecaster reports

We simulate a forecaster with no skill whatsoever: a fair coin, held for a run length drawn from the empirical distribution above. It is evaluated against real BANKNIFTY




Fig. 1. Autocorrelation of the forward return, overlapping versus non-overlapping, at two horizons. In both panels the overlapping series decays to zero at exactly lag k , which is the signature of arithmetic overlap rather than of serial structure.

TABLE III

A FORECASTER WITH EXACTLY ZERO SKILL, 30-MINUTE HORIZON, SCORED PER BAR VERSUS PER EVENT. BANKNIFTY; NIFTY 50 IS WITHIN 0.3 POINTS ON EVERY ROW.

	Per bar	Per k bars
Reported sample size per session	69	12
Stated binomial 95 % CI	± 11.8 pts	± 28.3 pts
Realised spread across sessions	± 19.6 pts	± 28.0 pts
Understatement factor	1.66$\times$	0.99$\times$
$P(\text{session} \geq 60\%)$	15.9 %	19.0 %
$P(\text{session} \geq 71\%)$	1.97 %	7.1 %
$P(\text{session} \geq 85\%)$	0.01 %	0.30 %

the stated interval matches the realised spread at $0.99\times$; scored per bar it is $1.66\times$ too narrow, because the reported n of 69 is fictitious. The effective sample is closer to 12. Second, the per-event column has *wider* absolute intervals and higher tail probabilities. This is not a contradiction: honest scoring correctly reports large uncertainty from a small sample, whereas per-bar scoring reports small uncertainty from a sample it does not have.

D. Panel multiplicity

A dashboard does not monitor one number. The system studied here displays two indices at four horizons, so eight hit-rates every session. Under independence,

$$P(\exists \text{ series} \geq 71\%) = 1 - (1 - 0.0197)^8 = 14.7\%,$$

or approximately one session in seven. A headline 71 % therefore arrives within the first fortnight of live operation with no edge whatsoever.

Our own 71 % was exactly this. Traced to source it was NIFTY 50, 30-minute horizon, on a single trending Friday: 202 of 285 per-minute forecasts scored as hits. The identical configuration scored 56 % the following Monday, and 51 % when re-evaluated out-of-sample at one observation per bar.

VI. A CORRECTED ESTIMATOR

Diagnosing a biased statistic is only half a contribution. This section supplies the closed form behind the distortion and an estimator that removes it, then validates the estimator by coverage rather than by argument.

A. Effective sample size

Let $h_t = \mathbf{1}\{\hat{y}_t = y_t\}$ denote the hit indicator of a forecaster issuing a k -bar-ahead call every bar, and let $\hat{p} = n^{-1} \sum_t h_t$ be the reported hit-rate. Practitioners quote $\text{Var}(\hat{p}) = p(1-p)/n$, which holds only for independent h_t . In general,

$$\text{Var}(\hat{p}) = \frac{p(1-p)}{n} \left[1 + 2 \sum_{j=1}^{n-1} \left(1 - \frac{j}{n} \right) \rho_j \right], \quad (1)$$

Zero-skill forecaster, scored every bar
(BANKNIFTY, 30m, 70,000 simulated sessions)




Fig. 2. Measured session hit-rate distribution of a zero-skill forecaster scored every bar. The multimodality is not noise: with roughly a dozen independent decisions per session the hit-rate lands on near-discrete values, which is the small effective sample made visible.

and NIFTY 50 bars and scored the way a live dashboard scores it. Results over approximately 70 000 simulated sessions per index appear in Table III and Fig. 2.

Two features of Table III deserve emphasis. First, the *calibration* column is the diagnostic: scored per event,

TABLE IV

COVERAGE OF A NOMINAL 95 % INTERVAL ON A SESSION HIT-RATE, AGAINST A FORECASTER OF KNOWN SKILL WITH REALISTIC PERSISTENCE. BANKNIFTY, 30-MINUTE HORIZON, NOMINAL $n = 69$; NIFTY 50 AGREES WITHIN 1 POINT THROUGHOUT.

Estimator	Coverage at skill			Width (pts)	n_{eff}
	50 %	55 %	60 %		
Naive binomial	39.6	39.5	38.9	21.4	69
HAC (Bartlett)	66.6	66.6	66.1	33.8	21
Block bootstrap	63.1	63.0	63.0	29.6	—
Run-structure	94.0	94.1	94.2	70.8	5.6

where $\rho_j = \text{Corr}(h_t, h_{t+j})$. The bracketed term is the variance inflation factor, and dividing n by it gives the effective sample size n_{eff} .

Two structures drive ρ_j , and the argument of Section V is that both are required. The outcome y_t is autocorrelated out to lag k because consecutive windows share bars, and the prediction $\hat{y}_t$ persists in runs. If predictions are redrawn each bar, h_t is i.i.d. whatever y_t does, every ρ_j vanishes, and $n_{\text{eff}} = n$.

Because the dependence here is block-shaped, (1) admits a direct form. If the sequence decomposes into runs of length L_i within which h_t is constant, then conditioning on the lengths gives $\text{Var}(\hat{p}) = p(1-p) \sum_i L_i^2 / n^2$, so

$$n_{\text{eff}} = \frac{n}{\sum_i L_i^2 / n} \xrightarrow{n \rightarrow \infty} \frac{n \mathbb{E}[L]}{\mathbb{E}[L^2]}. \quad (2)$$

The inflation is governed by $\mathbb{E}[L^2]$, not $\mathbb{E}[L]$. A run-length distribution with a heavy tail inflates the variance far beyond what its mean suggests, and that is precisely the regime a real forecaster occupies: on our production log the mean run is 15.2 minutes while the longest is 135.

B. Validation by coverage

An interval is judged by whether it covers. We generate a forecaster whose true skill is fixed by construction, with calls persisting according to the measured run-length distribution, and ask how often a nominal 95 % interval contains that known truth. Table IV reports four estimators.

The naive interval covers 40 % of the time while claiming 95 %. A standard HAC correction improves this substantially but still under-covers at 67 %, and a moving-block bootstrap reaches 63 %. Both fall short for the same reason: they mis-estimate how far the dependence runs. The conventional automatic Bartlett bandwidth, $4(n/100)^{2/9}$, evaluates to roughly three bars here, whereas a real forecaster holds a call for forty. Truncating the kernel at three bars discards most of the covariance that (1) is summing.

Estimating the inflation directly from the observed run structure removes that failure mode and restores coverage to 94 %.

C. What the corrected number says

The practical content is n_{eff} itself. A session that reports 69 forecasts carries roughly *six* independent observations,

TABLE V

MODEL-INDEPENDENT CEILING TEST, 30-MINUTE HORIZON, $n = 16\,856$ OUT-OF-SAMPLE PREDICTIONS PER INDEX.

Measurement	BANKNIFTY	NIFTY 50
Training accuracy	100.0 %	100.0 %
Out-of-sample, real labels	50.5 %	50.0 %
Out-of-sample, shuffled labels	50.2 %	50.2 %
1-NN oracle	48.8 %	50.2 %
Majority-class baseline	49.3 %	49.8 %
Positive control, $s = 0.10$	52.4 %	52.0 %
Positive control, $s = 0.30$	63.5 %	63.7 %

TABLE VI

OUT-OF-SAMPLE DIRECTIONAL ACCURACY, 30-MINUTE HORIZON, PURGED WALK-FORWARD, $n = 16\,856$ PER INDEX. Δ IS THE GAIN OVER THAT MODEL'S OWN SHUFFLED-LABEL CONTROL.

Model	BANKNIFTY	Δ	NIFTY 50
Majority-class baseline	49.3 %	—	49.8 %
Logistic regression	50.6 %	+0.8	49.3 %
Gradient boosting	50.7 %	+0.4	49.7 %
LightGBM	51.1 %*	+1.2	49.8 %
XGBoost	51.3 %**	+1.6	50.1 %
k -nearest neighbours	50.5 %	+0.2	50.1 %
LSTM	50.4 %	+0.9	49.4 %
Transformer	50.8 %*	+0.8	48.3 %
Temporal Fusion Transf.	49.4 %	+0.6	50.1 %

* $p < 0.05$; ** $p < 0.05/18$ (Bonferroni, all model-index tests).

and an honest interval on its hit-rate is about ± 35 points wide. Such an interval spans everything from incompetence to apparent mastery, which is the correct description of what one session of a per-minute forecast stream can establish.

The correction changes no point estimate. Only the interval moves. What it buys is an interval that means what it claims, and the discipline of reporting n_{eff} alongside n whenever forecasts overlap.

VII. PREDICTABILITY UNDER CORRECT SCORING

A. Is there anything to predict?

We first ask the question in a form independent of any model. An unconstrained random forest was fitted until it memorised the training data (Table V).

The model memorises perfectly and generalises at chance. Permuting the labels changes out-of-sample accuracy by 0.3 points on both indices. The positive controls are recovered at $p < 0.001$, establishing that the pipeline detects signal when signal is present.

B. Nine model families

Table VI reports nine families on identical splits, features, seeds and controls. Fig. 3 plots the same results with exact intervals.

We report the nuance honestly rather than rounding it to a clean null. Gradient-boosted models on BANKNIFTY do exhibit a small, statistically solid effect: XGBoost




Fig. 3. Nine model families with exact binomial intervals, each against its own shuffled-label control. The spread between architectures is smaller than the width of a single interval.

TABLE VII

STRATEGY SWEEP WITH MULTIPLICITY CONTROL, BANKNIFTY, 86 TRIALS, NON-OVERLAPPING 30-MINUTE HOLDS AT 4 BP COST.

Quantity	Value
Strategies evaluated	86
Positive Sharpe at all	2
Significantly profitable (one-sided)	0
Surviving Bonferroni	0
Surviving Benjamini–Hochberg, FDR 5%	0
Positive <i>and</i> stable in both halves	1
Best observed Sharpe	0.21
Bootstrap null, best-of-86, median	0.44
Bootstrap null, best-of-86, 95th pct.	1.31
$P(\text{null} \geq \text{observed})$	0.658

reaches 51.3% against a shuffled control of 49.7%, with $p = 0.001$ surviving Bonferroni correction across all eighteen model–index tests. The effect recurs at the 60-minute horizon, where LSTM reaches 51.2% ($p = 0.002$). It does *not* replicate on NIFTY 50 at any horizon. We take the honest reading to be a genuine but index-specific effect of roughly one percentage point. Section VIII shows that one point is nowhere near enough.

C. Eighty-six strategies and the statistics they require

Reporting the best of a large sweep is selection, not discovery. We evaluated 86 parameterised rules, covering momentum and mean reversion at ten lookbacks, six moving average crossover pairs, RSI, Bollinger, breakout and fade variants, MACD, gap, volatility regime, time-of-day, day-of-week and streak rules, on non-overlapping 30-minute holds with 4 bp round-trip cost (Table VII).

The best strategy in the sweep is *worse* than what selection alone typically manufactures. Two methodological notes:

One-sided tests. An earlier version of this analysis used a two-sided t -test and reported “42 of 86 significant.” Those were strategies that reliably *lose*. The hypothesis is profitability, so the test is one-sided; reliable losers are the cost of trading noise, not discoveries.

Where the standard corrections fail. The Deflated Sharpe Ratio [11] estimates its null from the cross-sectional variance of trial Sharpes, assuming a common-variance normal sample. Our library violates this badly: strategies differ by an order of magnitude in turnover, so their Sharpe dispersion is dominated by deterministic cost differences rather than by sampling noise, and the formula returns an implausible null benchmark of 7.77. We therefore replaced it with a circular block bootstrap [4], [5] that resamples forward returns in blocks while holding each strategy’s positions and turnover fixed, destroying the signal–outcome relationship and nothing else. Similarly, PBO via CSCV [12] is *degenerate* here. With nothing significantly profitable in-sample there is no selection to overfit, so we report it as such rather than as a result. Reporting a degenerate statistic as a finding would be an instance of the error this paper is about.

VIII. ECONOMIC SIGNIFICANCE

Directional accuracy is the field’s headline metric and is, alone, close to meaningless. Whether a given accuracy earns money depends on the size of the move relative to the cost of capturing it.

For a symmetric bracket with target and stop both at the typical move m and round-trip cost c , expectancy per trade is

$$\mathbb{E} = p(m - c) - (1 - p)(m + c) = m(2p - 1) - c,$$

so break-even requires

$$p^* = \frac{1}{2} + \frac{c}{2m}. \quad (3)$$




Fig. 4. Break-even accuracy against the best accuracy any model family reached, by horizon. Points beyond the dashed line at 100% are horizons at which no directional accuracy is sufficient.

TABLE VIII

BREAK-EVEN ACCURACY FROM MEASURED MOVE DISTRIBUTIONS AT 4 BP ROUND-TRIP COST. “ACHIEVED” IS THE BEST OF THE NINE MODEL FAMILIES.

Horizon	Med. move	c/m	p^*	Achieved
<i>BANKNIFTY</i>				
5 min	22.1	1.01	100.7 %	—
15 min	37.7	0.60	79.8 %	—
30 min	52.3	0.43	71.5 %	51.3 %
60 min	72.9	0.31	65.4 %	51.2 %
<i>NIFTY 50</i>				
5 min	8.1	1.22	111.2 %	—
15 min	14.1	0.70	84.9 %	—
30 min	19.7	0.50	75.1 %	50.1 %
60 min	27.7	0.35	67.7 %	—

Median move in index points.

The ratio c/m is therefore the entire problem. Table VIII and Fig. 4 evaluate (3) on measured move distributions.

Two consequences follow immediately.

At the shortest horizon the requirement exceeds 100 % on both indices. The median five-minute move is smaller than the round trip. There exists no accuracy, whether 90 % or 100 %, at which five-minute signal trading on these instruments is profitable at retail cost. The shorter the horizon, the more certainly it fails, which inverts the usual marketing of high-frequency retail strategies.

Instrument choice can make the requirement unreachable. Table IX evaluates (3) across cost regimes.

“Impossible” denotes $c > m$, hence $p^* > 1$: the round trip exceeds the typical move and a perfect forecaster loses on the median trade. The instrument most retail participants actually use, over the horizon they actually hold, has no directional accuracy that makes it profitable.

TABLE IX

BREAK-EVEN ACCURACY BY COST REGIME, BANKNIFTY 30-MINUTE HORIZON (MEDIAN MOVE 9.3 BP).

Cost	Instrument	p^*
1 bp	Exchange fees only, no slippage	55.4 %
2 bp	Tight futures, minimal slippage	60.8 %
4 bp	Realistic futures round trip	71.5 %
8 bp	Futures with adverse slippage	93.0 %
20 bp	ATM options: spread + 30 min theta	impossible
40 bp	OTM options or a wide book	impossible

A. Statistical significance is not economic significance

The BANKNIFTY effect of Section VII is real by every standard applied: $p = 0.001$ at $n = 16856$, Bonferroni-robust, replicated across horizons, and clearly separated from its shuffled control. It is also worthless, because 51.3 % against a 71.5 % requirement loses on every trade.

This distinction explains why fourteen months of accuracy optimisation produced no tradeable result even where it produced a statistically defensible one, and it is the reason we regard accuracy as the wrong headline metric for this task.

B. Reframing high-accuracy claims

Suppose a claimed 85 % accuracy at the 30-minute horizon were genuine. It would clear break-even by 13.5 points on BANKNIFTY futures and 9.9 on NIFTY 50, and would remain unprofitable on options. A claim of 85 % is therefore either an artefact of the kind measured in Section V, where a zero-skill system reaches that figure occasionally, or it is true and far less valuable than it sounds. Both readings argue against paying for it.

C. A filter that worked and still lost

A final corroboration. We trained a meta-label filter [13] in which a secondary model decides whether a primary

signal’s setup is worth taking, labelled on *points per minute* rather than on target-hit. The label choice matters: training on target-hit raises the hit-rate from 42 % to 53 % while making the money worse, because the model learns to prefer small, easy targets.

Out-of-sample over 106 unseen days on BANKNIFTY, the filter cut trades from 2.3 to 1.0 per day and raised quality from +1.44 to +1.82 points per minute, holding in both halves of the test. The total moved from −3777 to −981. A working filter applied to a system with no edge converts a large loss into a small one. It cannot manufacture a profit that is not in the underlying signal.

IX. WHAT IS PREDICTABLE INSTEAD

The fifteen production features are all transforms of one price series; reshaping price cannot reveal what price does not contain. Our positive results come from changing the target or the information source.

Volatility. A binary magnitude target, asking whether realised volatility over the next 30 minutes exceeds the training median, reaches 68 % out-of-sample on BANKNIFTY and 74 % on NIFTY 50 using an LSTM, stable across both test halves. This is consistent with volatility clustering being among the most robust properties of financial series [17], [18]. We attach the necessary caveat: a naive persistence rule already scores approximately 80 % on BANKNIFTY, so the genuine *edge over baseline* is on NIFTY 50 at the 30- and 60-minute horizons (+9 to +13 points). Reporting the raw 80 % without the baseline would repeat the error this paper documents.

Variance risk premium. Short at-the-money straddles held to expiry, priced on real end-of-day option premiums over two years, return +22.5 points per trade after 3 % costs at a 61 % win rate on NIFTY 50 over 349–509 trades. The effect is structural rather than predictive, being compensation for bearing variance risk [20], [21], and that is precisely why it survived a protocol which eliminated every predictive signal. It is also a short volatility payoff with a heavy left tail (worst trade −847 on NIFTY 50, −2002 on BANKNIFTY, where the same structure loses −7.3 points per trade), and must be implemented with defined risk. A high win rate here is a warning about tail exposure, not evidence of skill.

Option-chain positioning. Over 2181 days of end-of-day chains, four features survive after the price component is regressed out. They are the put–call ratio in open interest and in volume, its daily change, and distance from the max-pain strike, with information coefficients of 0.04–0.06 stable across both halves of a decade, on days where price-only features show nothing. The strongest is the *change* in put–call ratio rather than its level: flow beats position. An IC of 0.05 corresponds to roughly a 52 % hit-rate, which is real and, by (3), not tradeable on its own.

Edges decay. The single cost-positive signal we found was a contrarian rule on put–call ratio extremes: 204 trades, profit factor 1.11 over two years. It was already

dead when we discovered it. Directional accuracy of 74 % in 2024 fell to 36 %, with a profit factor of 0.74, over the final six months. This is the mechanism Timmermann and Granger describe [16], whereby discovered patterns self-destruct. It also bears on sustained-accuracy claims: a genuinely real edge in this data died inside two years, so a claim of stable accuracy over 76 trading days carries no information even if every number in it is true.

X. LIMITATIONS

We state these plainly because the paper’s argument is about honest reporting.

Two indices carry the primary intraday analysis, over seventeen months, in one market. The break-even model (3) assumes a symmetric bracket; asymmetric pay-offs lower the required hit-rate while also lowering the achieved one, and we did not sweep that trade-off. Our Temporal Fusion Transformer is a compact variant retaining variable selection, gated residual networks and interpretable attention, but omitting static covariates and the quantile head, and should not be read as the published architecture [28]. Some results in our research record come from scripts not retained; these are labelled in the repository and none are load-bearing for Sections V–VIII. Index spot carries no volume, so any VWAP-derived feature is misnamed, and whether a true volume-weighted average helps is untested. Finally, the order-flow hypothesis, where the literature locates genuine short-horizon alpha [22], [23], remains untested here for want of history: we run a live collector, but message-level history is not sold to retail accounts.

XI. CONCLUSION

The headline metric of intraday directional forecasting is not merely noisy. It is systematically biased upward by the way it is computed, and the bias is large enough to manufacture exactly the figures used in marketing: a forecaster with zero skill posts a ≥ 71 % session once in 51, and once every seven sessions across a modest panel, while reporting a confidence interval $1.66\times$ too narrow.

Scored correctly, nine model families reach at most 51.3 %, 86 strategies produce a winner worse than a bootstrap null, and the small effect that does survive every control sits twenty percentage points below its own break-even threshold. On the instruments and horizons most retail participants use, that threshold exceeds 100 %.

What remains is not direction. Magnitude is forecastable, the variance risk premium is real and structural, and order-flow data, which we could not obtain, is where the literature locates the remaining short-horizon alpha.

Reporting practice should change to match, and the change is cheap. Section VI gives an estimator whose interval covers at 94 % where the conventional one covers at 40 %, at no cost to the point estimate. Scoring unit, effective sample size, multiple-testing correction, a shuffled-label control and an explicit cost model are not optional

refinements to an accuracy figure. Without them the figure does not mean anything.

REPRODUCIBILITY

Code, the experiment harness, raw results, figure generation and the full research record, including five formally retracted claims, are available at <https://github.com/rachitranka25/accuracy-illusion>. Experiments use a fixed seed and pin native thread pools to one thread.

REFERENCES

- [1] L. P. Hansen and R. J. Hodrick, "Forward exchange rates as optimal predictors of future spot rates: An econometric analysis," *Journal of Political Economy*, vol. 88, no. 5, pp. 829–853, 1980.
- [2] M. Richardson and T. Smith, "Tests of financial models in the presence of overlapping observations," *The Review of Financial Studies*, vol. 4, no. 2, pp. 227–254, 1991.
- [3] W. K. Newey and K. D. West, "A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix," *Econometrica*, vol. 55, no. 3, pp. 703–708, 1987.
- [4] H. R. Künsch, "The jackknife and the bootstrap for general stationary observations," *The Annals of Statistics*, vol. 17, no. 3, pp. 1217–1241, 1989.
- [5] D. N. Politis and J. P. Romano, "The stationary bootstrap," *Journal of the American Statistical Association*, vol. 89, no. 428, pp. 1303–1313, 1994.
- [6] M. H. Pesaran and A. Timmermann, "A simple nonparametric test of predictive performance," *Journal of Business & Economic Statistics*, vol. 10, no. 4, 1992.
- [7] F. X. Diebold and R. S. Mariano, "Comparing predictive accuracy," *Journal of Business & Economic Statistics*, vol. 13, no. 3, pp. 253–263, 1995.
- [8] Y. Benjamini and Y. Hochberg, "Controlling the false discovery rate: A practical and powerful approach to multiple testing," *Journal of the Royal Statistical Society: Series B*, vol. 57, no. 1, pp. 289–300, 1995.
- [9] H. White, "A reality check for data snooping," *Econometrica*, vol. 68, no. 5, pp. 1097–1126, 2000.
- [10] C. R. Harvey, Y. Liu, and H. Zhu, "...and the cross-section of expected returns," *The Review of Financial Studies*, vol. 29, no. 1, pp. 5–68, 2016.
- [11] D. H. Bailey and M. López de Prado, "The deflated Sharpe ratio: Correcting for selection bias, backtest overfitting and non-normality," *The Journal of Portfolio Management*, vol. 40, no. 5, pp. 94–107, 2014.
- [12] D. H. Bailey, J. Borwein, M. López de Prado, and Q. J. Zhu, "The probability of backtest overfitting," *Journal of Computational Finance*, vol. 20, no. 4, pp. 39–70, 2017.
- [13] M. López de Prado, *Advances in Financial Machine Learning*. Hoboken, NJ, USA: Wiley, 2018.
- [14] E. F. Fama, "Efficient capital markets: A review of theory and empirical work," *The Journal of Finance*, vol. 25, no. 2, pp. 383–417, 1970.
- [15] A. W. Lo and A. C. MacKinlay, "Stock market prices do not follow random walks: Evidence from a simple specification test," *The Review of Financial Studies*, vol. 1, no. 1, pp. 41–66, 1988.
- [16] A. Timmermann and C. W. J. Granger, "Efficient market hypothesis and forecasting," *International Journal of Forecasting*, vol. 20, no. 1, pp. 15–27, 2004.
- [17] R. F. Engle, "Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation," *Econometrica*, vol. 50, no. 4, pp. 987–1007, 1982.
- [18] T. Bollerslev, "Generalized autoregressive conditional heteroskedasticity," *Journal of Econometrics*, vol. 31, no. 3, pp. 307–327, 1986.
- [19] A. R. Admati and P. Pfleiderer, "A theory of intraday patterns: Volume and price variability," *The Review of Financial Studies*, vol. 1, no. 1, pp. 3–40, 1988.
- [20] P. Carr and L. Wu, "Variance risk premiums," *The Review of Financial Studies*, vol. 22, no. 3, pp. 1311–1341, 2009.
- [21] G. Bakshi and N. Kapadia, "Delta-hedged gains and the negative market volatility risk premium," *The Review of Financial Studies*, vol. 16, no. 2, pp. 527–566, 2003.
- [22] P. N. Kolm, J. Turiel, and N. Westray, "Deep order flow imbalance: Extracting alpha at multiple horizons from the limit order book," *Mathematical Finance*, vol. 33, no. 4, pp. 1044–1081, 2023.
- [23] Z. Zhang, S. Zohren, and S. Roberts, "DeepLOB: Deep convolutional neural networks for limit order books," *IEEE Transactions on Signal Processing*, vol. 67, no. 11, pp. 3001–3012, 2019.
- [24] S. Gu, B. Kelly, and D. Xiu, "Empirical asset pricing via machine learning," *The Review of Financial Studies*, vol. 33, no. 5, pp. 2223–2273, 2020.
- [25] T. Fischer and C. Krauss, "Deep learning with long short-term memory networks for financial market predictions," *European Journal of Operational Research*, vol. 270, no. 2, pp. 654–669, 2018.
- [26] C. Krauss, X. A. Do, and N. Huck, "Deep neural networks, gradient-boosted trees, random forests: Statistical arbitrage on the S&P 500," *European Journal of Operational Research*, vol. 259, no. 2, pp. 689–702, 2017.
- [27] O. B. Sezer, M. U. Gudelek, and A. M. Ozbayoglu, "Financial time series forecasting with deep learning: A systematic literature review: 2005–2019," *Applied Soft Computing*, vol. 90, art. 106181, 2020.
- [28] B. Lim, S. Ö. Arik, N. Loeff, and T. Pfister, "Temporal fusion transformers for interpretable multi-horizon time series forecasting," *International Journal of Forecasting*, vol. 37, no. 4, pp. 1748–1764, 2021.
- [29] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [30] A. Vaswani, N. Shazeer, N. Parmar, J. Uszkoreit, L. Jones, A. N. Gomez, Ł. Kaiser, and I. Polosukhin, "Attention is all you need," in *Advances in Neural Information Processing Systems*, vol. 30, 2017, pp. 5998–6008.
- [31] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. 22nd ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining*, 2016, pp. 785–794.
- [32] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, "LightGBM: A highly efficient gradient boosting decision tree," in *Advances in Neural Information Processing Systems*, vol. 30, 2017, pp. 3146–3154.
- [33] L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [34] C. J. Clopper and E. S. Pearson, "The use of confidence or fiducial limits illustrated in the case of the binomial," *Biometrika*, vol. 26, no. 4, pp. 404–413, 1934.